



Conjoining Specifications

MARTÍN ABADI and LESLIE LAMPORT
Digital Equipment Corporation

We show how to specify components of concurrent systems. The specification of a system is the conjunction of its components' specifications. Properties of the system are proved by reasoning about its components. We consider both the decomposition of a given system into parts, and the composition of given parts to form a system.

Categories and Subject Descriptors: D.2.4 [Software Engineering]: Program Verification—*correctness proofs*; F.3.1 [Logics and Meanings of Programs]: Specifying and Verifying and Reasoning about Programs—*specification techniques*

General Terms: Theory, Verification

Additional Key Words and Phrases: Composition, concurrent programming, decomposition, liveness properties, modular specification, safety properties, temporal logic

1. INTRODUCTION

Large systems are built from smaller parts. We present a method for deducing properties of a system by reasoning about its components. We show how to represent an individual component Π_i by a formula S_i so that the parallel composition usually denoted **cobegin** $\Pi_1 \parallel \dots \parallel \Pi_n$ **coend** is represented by the formula $S_1 \wedge \dots \wedge S_n$. Composition is conjunction.

We reduce composition to conjunction not for the sake of elegance, but because it is the best way we know to prove properties of composite systems. Rigorous reasoning requires logic, and hence a language of logical formulas. It does not require a conventional programming language for describing systems. We find it most convenient to regard programs and circuit descriptions as low-level specifications, and to represent them in the same logic used for higher-level specifications. The logic we use is TLA, the Temporal Logic of Actions [Lamport 1994]. We do not discuss here the important problem of translating from a low-level TLA specification to an implementation in a conventional language.

The idea of representing concurrent programs and their specifications as formulas in a temporal logic was first proposed by Pnueli [1981]. It was later observed that, if specifications allow “stuttering” steps that leave the state unchanged, then $S_l \Rightarrow S_h$ asserts that S_l implements S_h [Lamport 1983]. Hence, proving that a lower-level specification implements a higher-level one was reduced to proving a formula in

Authors' address: Systems Research Center, Digital Equipment Corporation, 130 Lytton Avenue, Palo Alto, CA 94301.

Permission to make digital/hard copy of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage, the copyright notice, the title of the publication, and its date appear, and notice is given that copying is by permission of ACM, Inc. To copy otherwise, to republish, to post on servers, or to redistribute to lists, requires prior specific permission and/or a fee.

© 1995 ACM 0164-0925/95/0500-0507\$03.50

the logic. Still later, it was noticed that the formula $\exists x : S$ specifies the same system as S except with the variable x hidden [Abadi and Lamport 1991; Lamport 1989], and variable hiding became logical quantification. The idea of composition as conjunction has also been suggested [Abadi and Plotkin 1993; Abramsky and Jagadeesan 1994; Zave and Jackson 1993], but our method for reducing composition to conjunction is new.

To deduce useful properties of a component, we must specify its environment. No component will exhibit its intended behavior in the presence of a sufficiently hostile environment. For example, a combinational circuit will not produce an output in the intended range if some input line, instead of having a 0 or a 1, has an improper voltage level of $1/2$. The specification of the circuit's environment must rule out such improper inputs.

How we reason about a composite system depends on how it was formed. Composite specifications arise in two ways: by *decomposing* a given system into smaller parts and by *composing* given parts to form a larger system. These two situations call for two methods of writing component specifications that differ in their treatment of the environment. This difference leads in turn to different proof rules.

When decomposing a specification, the environment of each component is assumed to be the other components, and is usually left implicit. To reason about a component, we must state what we are assuming about its environment, and then prove that this assumption is satisfied by the other components. The Decomposition Theorem of Section 4 provides the needed proof rule. It reduces the verification of a complex, low-level system to proving properties of a higher-level specification and properties of one low-level component at a time. Decomposing proofs in this way allows us to apply decision procedures to verifications that hitherto required completely hand-guided proofs [Kurshan and Lamport 1993].

When specifying a reusable component, without knowing precisely where it will be used, we must make explicit what it assumes of its environment. We therefore assert that the component satisfies a guarantee M only as long as its environment satisfies an assumption E . This assumption/guarantee property [Jones 1983] is denoted $E \dashv M$. To show that a composition of reusable components satisfies a specification S , we must prove a formula of the form $(E_1 \dashv M_1) \wedge \dots \wedge (E_n \dashv M_n) \Rightarrow S$, where S may again be an assumption/guarantee property. We prove such a formula with the Composition Theorem of Section 5. This theorem allows us to reason about assumption/guarantee specifications using well-established, effective methods for reasoning about specifications of complete systems.

In the following section, we examine the issues that arise in decomposition and composition. Our discussion is informal, because we wish to show that these issues are fundamental, not artifacts of a particular formalism. We treat these topics formally in Sections 4 and 5. Section 3 covers the formal preliminaries. A comparison with related work appears in the conclusion. Proofs are relegated to the appendix.

2. AN INFORMAL OVERVIEW

2.1 Decomposing Complete Systems

A complete system is one that is self-contained; it may be observed, but it does not interact with the observer. A program is a complete system, provided we model

ACM Transactions on Programming Languages and Systems, Vol. 17, No. 3, May 1995.

inputs as being generated nondeterministically by the program itself.

As a tiny example of a complete system, we consider a program for computing a GCD (greatest common divisor), for which we have devised an informal programming-language notation. Statements within angle brackets are executed atomically; **loop-endloop** keywords enclose an infinite loop; **cobegin-coend** keywords enclose parallel statements, separated by \parallel ; and semicolon has its usual meaning. When writing processes, we will also mark variables as output, input, or internal; a process cannot change its input variables or access the internal variables of another process.

```
PROGRAM GCD
  var a initially 233344, b initially 233577899;
  cobegin loop < if a > b then a := a - b > endloop
    ||
    loop < if b > a then b := b - a > endloop coend
```

Program GCD satisfies the correctness property that eventually a and b become and remain equal to the gcd of 233344 and 233577899. We make no distinction between programs and properties, writing them all as TLA formulas. If formula M_{gcd} represents program GCD, and formula P_{gcd} represents the correctness property, then the program implements the property iff (if and only if) M_{gcd} implies P_{gcd} . Thus, correctness of program GCD is verified by proving $M_{gcd} \Rightarrow P_{gcd}$.

In hierarchical development, one decomposes the specification of a system into specifications of its parts. As explained in Section 4, the specification M_{gcd} of program GCD can be written as $M_a \wedge M_b$, where M_a asserts that a initially equals 233344 and is repeatedly decremented by the value of b whenever $a > b$, and where M_b is analogous. The formulas M_a and M_b are the specifications of two processes Π_a and Π_b . We can write Π_a and Π_b as

PROCESS Π_a <pre>output var a initially 233344; input var b; loop < if a > b then a := a - b > endloop</pre>	PROCESS Π_b <pre>output var b initially 233577899; input var a; loop < if b > a then b := b - a > endloop</pre>
---	--

One decomposes a specification in order to refine the components separately. We can refine the GCD program, to remove simultaneous atomic accesses to both a and b , by refining process Π_a to

```
PROCESS  $\Pi_a^l$ 
  output var a initially 233344;
  internal var ai;
  input var b;
  loop < ai := b >; if < a > ai > then < a := a - ai > endloop
```

and refining Π_b to the analogous process Π_b^l .

The composition of processes Π_a^l and Π_b^l correctly implements program GCD. This is expressed in TLA by the assertion that $M_a^l \wedge M_b^l$ implies $M_a \wedge M_b$, where M_a^l and M_b^l are the formulas representing Π_a^l and Π_b^l .

We would like to decompose the proof of $M_a^l \wedge M_b^l \Rightarrow M_a \wedge M_b$ into proofs of $M_a^l \Rightarrow M_a$ and $M_b^l \Rightarrow M_b$. These proofs would show that Π_a^l implements Π_a and that Π_b^l implements Π_b .

	<i>initial</i>	<u>37</u>	<u>37</u>	<u>4</u>	<u>4</u>	<u>19</u>	
	<i>state</i>	<u>sent</u>	<u>acked</u>	<u>sent</u>	<u>acked</u>	<u>sent</u>	
<i>c.ack:</i>	0	0	1	1	0	0	...
<i>c.sig:</i>	0	1	1	0	0	1	...
<i>c.val:</i>	—	37	37	4	4	19	...

Fig. 1. The two-phase handshake protocol for a channel *c*.

Unfortunately, Π_a^l does not implement Π_a because, in the absence of assumptions about when its input b can change, Π_a^l can behave in ways that process Π_a cannot. Process Π_a can decrement a only by the current value of b , but Π_a^l can decrement a by a previous value of b if b changes between the assignment to ai and the assignment to a . Similarly, Π_b^l does not implement Π_b .

Process Π_a^l does correctly implement process Π_a in a context in which b does not change when $a > b$. This is expressed in TLA by the formula $E_a \wedge M_a^l \Rightarrow M_a$, where E_a asserts that b does not change when $a > b$. Similarly, $E_b \wedge M_b^l \Rightarrow M_b$ holds, for the analogous E_b . The Decomposition Theorem of Section 4.3 allows us to deduce $M_a^l \wedge M_b^l \Rightarrow M_a \wedge M_b$ from approximately the following hypotheses:

$$\begin{aligned} E_a \wedge M_a^l &\Rightarrow M_a \\ E_b \wedge M_b^l &\Rightarrow M_b \\ M_a \wedge M_b &\Rightarrow E_a \wedge E_b \end{aligned} \tag{1}$$

The third hypothesis holds because the composition of processes Π_a and Π_b does not allow a to change when $b > a$ or b to change when $a > b$.

Observe that E_a asserts only the property of Π_b^l needed to guarantee that Π_a^l implements Π_a . In a more complicated example, E_a will be significantly simpler than M_b^l , the full specification of Π_b^l . Verifying these hypotheses will therefore be easier than proving $M_a^l \wedge M_b^l \Rightarrow M_a \wedge M_b$ directly, since this proof requires reasoning about the specification $M_a^l \wedge M_b^l$ of the complete low-level program.

One cannot really deduce $M_a^l \wedge M_b^l \Rightarrow M_a \wedge M_b$ from the hypotheses (1). For example, (1) is trivially satisfied if E_a , E_b , M_a , and M_b all equal false; but we cannot deduce $M_a^l \wedge M_b^l \Rightarrow \text{false}$ for arbitrary M_a^l and M_b^l . The precise hypotheses of the Decomposition Theorem are more complicated, and we must develop a number of formal concepts in order to state them. We also develop results that allow us to discharge these more complicated hypotheses by proving conditions essentially as simple as (1).

2.2 Composing Open Systems

An open system is one that interacts with an environment it does not control. In our examples, we consider systems that communicate by using a standard two-phase handshake protocol [Mead and Conway 1980] to send values over channels. The state of a channel *c* is described by three components: the value *c.val* that is being sent, and two bits *c.sig* and *c.ack* used for synchronization. We let *c.snd* denote the pair $\langle c.sig, c.val \rangle$. Figure 1 shows the sequence of states assumed in sending the sequence of values 37, 4, 19, The channel is ready to send when *c.sig = c.ack*. A value *v* is sent by setting *c.val* to *v* and complementing *c.sig*. Receipt of the value is acknowledged by complementing *c.ack*.

We consider an *N*-element queue with input channel *i* and output channel *o*. It

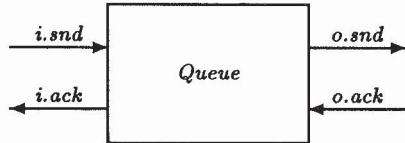


Fig. 2. A queue.

```

PROCESS QUEUE
  output var i.ack, o.sig initially 0,
        o.val;
  internal var q initially <>;
  input var i.sig, i.val, o.ack;
  cobegin
    loop < if (i.ack ≠ i.sig) ∧ (|q| < N)
           then q := q ∘ (i.val);
           i.ack := 1 - i.ack > endloop
    ||
    loop < if (o.ack = o.sig) ∧ (|q| > 0)
           then o.val := head(q);
           q := tail(q);
           o.sig := 1 - o.sig > endloop
  coend

```

Fig. 3. A queue process.

is depicted in Figure 2. To describe the queue, we use the programming-language constructs introduced in Section 2.1; in particular, we write large atomic actions within angle brackets. We also introduce the following notation for finite sequences: $|\rho|$ denotes the length of sequence ρ , which equals 0 if ρ is empty; $Head(\rho)$ and $Tail(\rho)$ as usual denote the head (first element) and the tail of sequence ρ , if ρ is nonempty; and $\rho \circ \tau$ denotes the concatenation of sequences ρ and τ . Moreover, angle brackets are used to form sequences; so $\langle \rangle$ denotes the empty sequence, and $\langle e \rangle$ denotes the sequence with e as its only element. With this notation, the queue can be written as in Figure 3.

Let QM be the TLA formula that represents this queue process. It might seem natural to take QM as the specification of the queue. However, this specification would be difficult or impossible to implement because it states that the queue behaves properly even if the environment does not obey the communication protocol. For example, in a lower-level implementation, reading the input $o.ack$ and setting the outputs $o.sig$ and $o.val$ would be separate actions. If the environment changed $o.ack$ between these actions, the implementation could violate the requirement that it change $o.val$ only when $o.ack = o.sig$. This problem is not an artifact of our particular representation of the queue; actual hardware implementations of a queue can enter metastable states, consequently producing bizarre, unpredictable behavior, if their inputs are changed when they are not supposed to be [Mead and Conway 1980].

A specification of the queue should allow executions in which the queue performs correctly; it should not rule out bad behavior of the queue caused by the environment performing incorrectly. Such a specification can be written in the assumption/guarantee style, a generalization of the traditional pre/postcondition style for sequential programs. An assumption/guarantee specification asserts that the system provides a guarantee M if its environment satisfies an assumption E . For the queue, M is the formula QM , and E asserts that the environment obeys

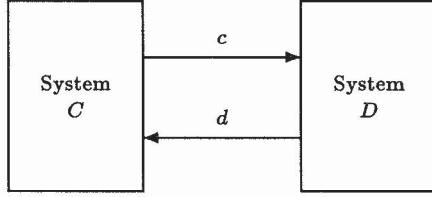


Fig. 4. A simple example.

the communication protocol.

It is not obvious how to reason about the composition of systems described by assumption/guarantee specifications. The basic problem is illustrated by the simple case of two systems, one guaranteeing M_c assuming M_d , and the other guaranteeing M_d assuming M_c . Since each system guarantees to satisfy the other's environment assumption, we would like to conclude that their composition implements the specification $M_c \wedge M_d$ unconditionally, with no environment assumption. Can we? We attempt to answer this question by considering two simple examples, based on Figure 4.

In the first example:

- M_c^0 asserts that c always equals 0.
- M_d^0 asserts that d always equals 0.

We can implement these specifications with the following two processes.

PROCESS Π_c

```

output var c initially 0;
input var d;
loop <c := d> endloop
  
```

PROCESS Π_d

```

output var d initially 0;
input var c;
loop <d := c> endloop
  
```

Process Π_c guarantees M_c^0 assuming M_d^0 , and process Π_d guarantees M_d^0 assuming M_c^0 . Clearly, their composition leaves c and d unchanged, so it implements $M_c^0 \wedge M_d^0$.

In the second example:

- M_c^1 asserts that c eventually equals 1.
- M_d^1 asserts that d eventually equals 1.

The same processes Π_c and Π_d implement the specifications in this case too; process Π_c guarantees M_c^1 assuming M_d^1 , and process Π_d guarantees M_d^1 assuming M_c^1 . However, since their composition leaves c and d unchanged, it does not implement $M_c^1 \wedge M_d^1$.

Our conclusion in the first example does not depend on the particular choice of processes Π_c and Π_d . We can deduce directly from the assumption/guarantee specifications that the composition must implement $M_c^0 \wedge M_d^0$, because the first process to change its output variable would violate its guarantee before its assumption had been violated. This argument does not apply to the second example, because violating M_c^1 and M_d^1 are sins of omission that do not occur at any particular instant. A property that can be made false only by being violated at some instant is called a safety property [Alpern and Schneider 1985]. As the examples suggest, reasoning about the composition of assumption/guarantee specifications is easiest when assumptions are safety properties.

The argument that the composition should implement $M_c^0 \wedge M_d^0$ in the first example rests on the requirement that a process maintains its guarantee until after

the environment violates its assumption. In other words, we interpret the assumption/guarantee specification as an assertion that the guarantee M can become false only after the assumption E becomes false. We write this assertion as the formula $E \dashv M$. Section 5 discusses this form of specification.

Our rules for reasoning about the composition of assumption/guarantee specifications are embodied in the Composition Theorem of Section 5.2. With the Composition Theorem, we can prove that the conjunction of the assumption/guarantee specifications $M_c^0 \dashv M_d^0$ and $M_d^0 \dashv M_c^0$ implies $M_c^0 \wedge M_d^0$. We can also prove more substantial results—for example, that the composition of queues implements a larger queue. Verifying the hypotheses of the theorem requires reasoning only about complete systems, so the theorem allows us to handle assumption/guarantee specifications as easily as complete-system specifications.

3. PRELIMINARIES

3.1 TLA

3.1.1 Review of the Syntax and Semantics. A state is an assignment of values to variables. (Technically, our variables are the “flexible” variables of temporal logic that correspond to the variables of programming languages; they are distinct from the variables of first-order logic.) A behavior is an infinite sequence of states. Semantically, a TLA formula F is true or false of a behavior; we say that F is *valid*, and write $\models F$, iff it is true of every behavior. Syntactically, TLA formulas are built up from state functions using Boolean operators (\neg , \wedge , \vee , \Rightarrow [implication], and $=$ [equivalence]) and the operators $', \square$, and \exists , as described below.

A *state function* is like an expression in a programming language. Semantically, it assigns a value to each state—for example $3 + x$ assigns to state s three plus the value of the variable x in s . A *state predicate* is a Boolean-valued state function. An *action* is a Boolean-valued expression containing primed and unprimed variables. Semantically, an action is true or false of a pair of states, with primed variables referring to the second state—for example, $x + 1 > y'$ is true for $\langle s, t \rangle$ iff the value of $x + 1$ in s is greater than the value of y in t . A pair of states satisfying action \mathcal{A} is called an \mathcal{A} *step*. We say that \mathcal{A} is *enabled* in state s iff there exists a state t such that $\langle s, t \rangle$ is an \mathcal{A} step—for example, $(x > 0) \wedge (x + 1 > y')$ is enabled only in states where $x > 0$. The state predicate *Enabled* \mathcal{A} is true for state s iff \mathcal{A} is enabled in s . We write v' for the expression obtained by priming all the variables of the state function v , and $[\mathcal{A}]_v$ for $\mathcal{A} \vee (v' = v)$, so an $[\mathcal{A}]_v$ step is either an \mathcal{A} step or a step that leaves v unchanged.

As usual in temporal logic, if F is a formula then $\square F$ is a formula that means that F is always true, and $\diamond F$, an abbreviation for $\neg \square \neg F$, means that F is eventually true. In addition, if \mathcal{A} is an action and v is a state function then $\square [\mathcal{A}]_v$ is a formula; $\diamond [\mathcal{A}]_v$ is an abbreviation for $\neg \square \neg [\mathcal{A}]_v$. Using \square and “enabled” predicates, we can define fairness operators WF and SF. The *weak-fairness* formula $WF_v(\mathcal{A})$ asserts of a behavior that either there are infinitely many \mathcal{A} steps that change v , or there are infinitely many states in which such steps are not enabled. This can be written $(\square \diamond [\mathcal{A}]_v) \vee (\square \diamond \neg \text{Enabled } [\mathcal{A}]_v)$. The *strong-fairness* formula $SF_v(\mathcal{A})$ asserts that either there are infinitely many \mathcal{A} steps that change v , or there are only finitely many states in which such steps are enabled. This can be written

$$(\square \diamond (\mathcal{A})_v) \vee (\diamond \square \neg \text{Enabled } (\mathcal{A})_v).$$

The formula $\exists x : F$ means essentially that there is some way of choosing a sequence of values for x such that the temporal formula F holds. We think of $\exists x : F$ as “ F with x hidden” and call x an internal variable of $\exists x : F$. Both x and x' are bound by $\exists x$ in $\exists x : F$. If x is a tuple of variables $\langle x_1, \dots, x_k \rangle$, we write $\exists x : F$ for $\exists x_1 : \dots \exists x_k : F$.

The standard way of specifying a system in TLA is with a formula in the “canonical form” $\exists x : \text{Init} \wedge \square[\mathcal{N}]_v \wedge L$, where Init is a predicate and L a conjunction of fairness conditions. This formula asserts that there exists a sequence of values for x such that (1) Init is true for the initial state, (2) every step of the behavior is an \mathcal{N} step or leaves the state function v unchanged, and (3) L holds. For example, the specification M_{gcd} of the complete high-level GCD program is written in canonical form by taking¹

$$\begin{aligned} \text{Init} &\triangleq (a = 233344) \wedge (b = 233577899) \\ \mathcal{N} &\triangleq \begin{aligned} &\vee (a > b) \wedge (a' = a - b) \wedge (b' = b) \\ &\vee (b > a) \wedge (b' = b - a) \wedge (a' = a) \end{aligned} \\ v &\triangleq \langle a, b \rangle \\ L &\triangleq \text{WF}_v(\mathcal{N}) \end{aligned} \tag{2}$$

Intuitively, a variable represents some part of the universe, and a behavior represents a possible complete history of the universe. A system Π is represented by a TLA formula M that is true for precisely those behaviors that represent histories in which Π is running. We make no formal distinction between systems, specifications, and properties; they are all represented by TLA formulas, which we usually call specifications.

3.1.2 Interleaving and Noninterleaving Representations. Let ξ and ψ be two objects, represented by the variables x and y , respectively. When representing a history of the universe as a behavior, we can describe concurrent changes to ξ and ψ either by a single simultaneous change to x and y , or by separate changes to x and y in some order. If the changes to ξ and ψ are directly linked, then it is usually most convenient to describe their concurrent change by a single change to both x and y . However, if the changes are independent, then we are free to choose whether or not to allow simultaneous changes to x and y . An *interleaving* representation is one in which such simultaneous changes are disallowed.

When changes to ξ and ψ are directly linked, we often think of x and y as output variables of a single component. An interleaving representation is then one in which simultaneous changes to output variables of different processes are disallowed. The absence of such simultaneous changes can be expressed as a TLA formula. For a system with n components in which v_i is the tuple of output variables of component i , interleaving is expressed by the formula

$$\text{Disjoint}(v_1, \dots, v_n) \triangleq \bigwedge_{i \neq j} \square[(v'_i = v_i) \vee (v'_j = v_j)]_{(v_i, v_j)}$$

¹We let a list of formulas bulleted with \wedge or \vee denote the conjunction or disjunction of the formulas, using indentation to eliminate parentheses. We also let \Rightarrow have lower precedence than the other Boolean operators.

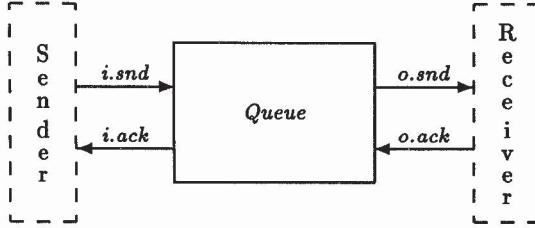


Fig. 5. The complete system of queue plus environment.

We have found that, in TLA, interleaving representations are usually easier to write and to reason about. Moreover, an interleaving representation is adequate for reasoning about a system if the system is modeled at a sufficiently fine grain of atomicity. However, as discussed below, TLA also works for noninterleaving representations. TLA does not mandate any particular method for representing systems. Indeed, one can write specifications that are intermediate between interleaving and noninterleaving representations.

3.1.3 The Queue Example. We now give a TLA specification of the queue of natural numbers of length N , which was described informally in Section 2.2 and illustrated in Figure 2. As in Section 2.2, we write $c.snd$ for the pair $\langle c.sig, c.val \rangle$ for a channel c ; we also write c for the triple $\langle c.sig, c.ack, c.val \rangle$.

A channel is initially ready for sending, so the initial condition on wire c is the predicate $CInit(c)$ defined by

$$CInit(c) \triangleq (c.sig = c.ack = 0)$$

The operations of sending a value v and acknowledging receipt of a value on channel c are represented by the following $Send(v, c)$ and $Ack(c)$ actions.

$$\begin{array}{ll} Send(v, c) \triangleq & \wedge c.sig = c.ack \\ & \wedge c.snd' = \langle 1 - c.sig, v \rangle \\ & \wedge c.ack' = c.ack \\ Ack(c) \triangleq & \wedge c.sig \neq c.ack \\ & \wedge c.ack' = 1 - c.ack \\ & \wedge c.snd' = c.snd \end{array}$$

To represent the queue as a complete system, we add an environment that sends arbitrary natural numbers over channel i and acknowledges receipt of values on channel o . The resulting complete system is shown in Figure 5.

The TLA formula CQ specifying the queue is defined in Figure 6. It has the canonical form $\exists x : Init \wedge \square[\mathcal{N}]_v \wedge L$, where:

- x is the internal variable q , which represents the sequence of values received on the input channel i but not yet sent on the output channel o .
- $Init$ is written as the conjunction $Init_E \wedge Init_M$ of initial predicates for the environment and component. (We arbitrarily consider the initial conditions on a channel to be part of the sender's initial predicate.)
- \mathcal{N} is the disjunction of two actions: \mathcal{Q}_M , describing the steps taken by the component, and $\mathcal{Q}_E \wedge (q' = q)$, describing steps taken by the environment (which leave q unchanged). Action \mathcal{Q}_M is the disjunction of actions Enq and Deq . An Enq step acknowledges receipt of a value on i and appends the value to q ; it is enabled only when q has fewer than N elements. A Deq step removes the first element of q and sends it on o . Action \mathcal{Q}_E is the disjunction of Put , which

$Init_E$	\triangleq	$CInit(i)$	Environment Actions
Put	\triangleq	$(\exists v \in \text{Nat} : Send(v, i)) \wedge (o' = o)$	
Get	\triangleq	$Ack(o) \wedge (i' = i)$	
Q_E	\triangleq	$Get \vee Put$	
$Init_M$	\triangleq	$CInit(o) \wedge (q = \langle \rangle)$	Component Actions
Enq	\triangleq	$\wedge q < N$ $\wedge Ack(i) \wedge (q' = q \circ \langle i.val \rangle)$ $\wedge o' = o$	
Deq	\triangleq	$\wedge q > 0$ $\wedge Send(Head(q), o) \wedge (q' = Tail(q))$ $\wedge i' = i$	
Q_M	\triangleq	$Enq \vee Deq$	
ICL	\triangleq	$WF_{\langle i, o, q \rangle}(Q_M)$	Complete-System Specification
ICQ	\triangleq	$\wedge Init_E \wedge Init_M$ $\wedge \square \left[\begin{array}{l} \vee Q_E \wedge (q' = q) \\ \vee Q_M \end{array} \right]_{\langle i, o, q \rangle}$ $\wedge ICL$	
CQ	\triangleq	$\exists q : ICQ$	

Fig. 6. The specification CQ of the complete queue. (Formulas $CInit$, $Send$, and Ack are defined in the text.)

sends an arbitrary number on channel i , and Get , which acknowledges receipt of a number on channel o .

— v is the tuple $\langle i, o, q \rangle$ of all relevant variables.²

— L is the weak-fairness condition ICL , which is defined to be $WF_{\langle i, o, q \rangle}(Q_M)$, and asserts that a component step cannot remain forever possible without occurring. It can be shown that a logically equivalent specification is obtained if this condition is replaced with $WF_{\langle i, o, q \rangle}(Enq) \wedge WF_{\langle i, o, q \rangle}(Deq)$.

Formula CQ gives an interleaving representation of a queue; simultaneous steps by the queue and its environment are not allowed. Moreover, simultaneous changes to the two inputs $i.snd$ and $o.ack$ are disallowed, as are simultaneous changes to the two outputs $i.ack$ and $o.snd$. In Section 4, we describe a noninterleaving representation of the queue.

3.2 Implementation

A specification M^l implies a specification M iff every behavior that satisfies M^l also satisfies M ; hence proving $M^l \Rightarrow M$ shows that the system Π^l represented by M^l implements the system or property Π represented by M . Note that if M^l is inconsistent (equivalent to false), then $M^l \Rightarrow M$ holds vacuously, but an inconsistent M^l does not represent any system Π^l .

The formula $M^l \Rightarrow M$ is proved by applying a handful of simple rules [Lamport 1994]. When M has the form $\exists x : \widehat{M}$, a key step in the proof is finding a *refinement*

²Informally, we write $\langle i, o, q \rangle$ for the concatenation of the tuples i , o , and $\langle q \rangle$.

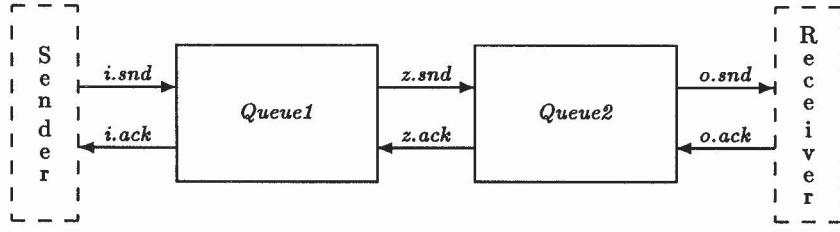


Fig. 7. A complete system containing two queues in series.

$$\begin{aligned}
 ICDQ &\triangleq \wedge Init_E \wedge Init_M^{[1]} \wedge Init_M^{[2]} \\
 &\quad \wedge \square \left[\begin{array}{l} \vee Q_E \wedge \langle q_1, q_2, z \rangle' = \langle q_1, q_2, z \rangle \\ \vee Q_M^{[1]} \wedge \langle q_2, o \rangle' = \langle q_2, o \rangle \\ \vee Q_M^{[2]} \wedge \langle q_1, i \rangle' = \langle q_1, i \rangle \end{array} \right]_{\langle i, o, z, q_1, q_2 \rangle} \\
 &\quad \wedge ICL^{[1]} \wedge ICL^{[2]} \\
 CDQ &\triangleq \exists q_1, q_2 : ICDQ
 \end{aligned}$$

Fig. 8. Specification of the complete double-queue system of Figure 7.

mapping—a tuple of state functions \bar{x} such that M^l implies \widehat{M} , where \widehat{M} is the formula obtained by substituting \bar{x} for x (and therefore $(\bar{x})'$ for x') in \widehat{M} . Under reasonable assumptions, such a refinement mapping exists when $M^l \Rightarrow \exists x : \widehat{M}$ is valid [Abadi and Lamport 1991].

As an example, we show that the system composed of two queues in series, shown in Figure 7, implements a single larger queue. We first specify the composite queue. Let $F[e_1/v_1, \dots, e_n/v_n]$ denote the result of (simultaneously) substituting each expression e_i for v_i in a formula F . For example, if Get is defined as in Figure 6, then $Get[z/i]$ equals $Ack(o) \wedge (z' = z)$. For any formula F , let

$$F^{[1]} \triangleq F[z/o, q_1/q] \quad F^{[2]} \triangleq F[z/i, q_2/q]$$

In Figure 8, the specification CDQ of the complete system, consisting of the double queue and its environment, is defined in terms of the formulas from Figure 6. We think of the complete system as containing three components: the environment and the two queues. The initial condition is the conjunction of the initial conditions of each component. The next-state action consists of three disjuncts, representing actions of each of the three components that leave other components' variables unchanged. Finally, we take as the liveness condition the conjunction of the fairness conditions of the two queues.

We now show that the composite queue implements a $(2N + 1)$ -element queue. (The “+1” arises because the internal channel z acts as a buffer element.) The correctness condition is $CDQ \Rightarrow CQ^{[\text{dbl}]}$, where $F^{[\text{dbl}]}$ denotes $F[(2N + 1)/N]$, for any formula F . This is proved by showing $ICDQ \Rightarrow \overline{ICQ^{[\text{dbl}]}}$, with the refinement

mapping defined by

$$\bar{q} \triangleq \begin{array}{ll} \text{if } z.sig = z.ack & \text{then } q_1 \circ q_2 \\ & \text{else } q_1 \circ \langle z.val \rangle \circ q_2 \end{array}$$

The formula $ICDQ \Rightarrow \overline{ICQ^{[dbl]}}$ can be proved by standard TLA reasoning of the kind described by Lamport [1994].

3.3 Conditional Implementation

Instead of proving that a specification M^l implements a specification M , we sometimes want to prove the weaker condition that M^l implements M assuming a formula G . In other words, we want to prove $G \Rightarrow (M^l \Rightarrow M)$, which is equivalent to $G \wedge M^l \Rightarrow M$. The formula G may express one or more of the following:

- A law of nature. For example, in a real-time specification, G might assert that time increases monotonically. If the current time is represented by the variable now , this assumption is expressed by the formula $(now \in \mathbf{R}) \wedge \square[now' \in (now, \infty)]_{now}$, where \mathbf{R} is the set of real numbers.
- An interface refinement, where G expresses the relation between a low-level tuple l of variables and its high-level representation as a tuple h of variables. For example, l might be a low-level interface representing the transmission of sequences of bits over a wire, and h could be the high-level interface in which the sending of seven successive bits is interpreted as the transmission of a single ASCII character.
- An assumption about how reality is translated into the formalism of behaviors. In particular, G may assert an interleaving assumption—for example, an assumption of the form $Disjoint(v_1, \dots, v_n)$.

Conditional implementation, with an explicit formula G , is needed only for open systems. For a complete system, the properties expressed by G can easily be made part of the system specification. For example, the system can include a component that advances time. In contrast, it can be difficult to include G in the specification of an open system,

3.4 Safety and Closure

3.4.1 Definition of Closure. A finite sequence of states is called a finite behavior. For any formula F and finite behavior ρ , we say that ρ satisfies F iff ρ can be extended to an infinite behavior that satisfies F . For convenience, we say that the empty sequence $\langle \rangle$ satisfies every formula (even false).

A *safety property* is a formula that is satisfied by an infinite behavior σ iff it is satisfied by every prefix of σ [Alpern and Schneider 1985]. For any predicate $Init$, action \mathcal{N} , and state function v , the formula $Init \wedge \square[\mathcal{N}]_v$ is a safety property. It can be shown that, for any TLA formula F , there is a TLA formula $C(F)$, called the *closure of F* , such that a behavior σ satisfies $C(F)$ iff every prefix of σ satisfies F . Formula $C(F)$ is the strongest safety property such that $\models F \Rightarrow C(F)$. Proposition 1 below implies that $C(Init \wedge \square[\mathcal{N}]_v \wedge L)$ equals $Init \wedge \square[\mathcal{N}]_v$, when L is the conjunction of suitable fairness properties.

3.4.2 Machine Closure. When writing a specification in the form $\text{Init} \wedge \square[\mathcal{N}]_v \wedge L$, we expect L to constrain infinite behaviors, not finite ones. Formally, this means that the closure of $\text{Init} \wedge \square[\mathcal{N}]_v \wedge L$ should be $\text{Init} \wedge \square[\mathcal{N}]_v$. A pair of properties (P, L) is called *machine closed* iff $\mathcal{C}(P \wedge L)$ equals P [Abadi and Lamport 1991]. (We often say informally that $P \wedge L$ is machine closed.)

Proposition 1 below, which we have already proved [Abadi and Lamport 1994], shows that we can use fairness properties to write machine-closed specifications. The proposition relies on the following definition: an action \mathcal{A} is a *subaction* of a safety property P iff for every finite behavior $\rho = \langle r_0, \dots, r_n \rangle$, if ρ satisfies P and \mathcal{A} is enabled in state r_n , then there exists a state r_{n+1} such that $\langle r_0, \dots, r_{n+1} \rangle$ satisfies P and $\langle r_n, r_{n+1} \rangle$ is an \mathcal{A} step. It follows from this definition of subaction that, if \mathcal{A} implies \mathcal{N} , then \mathcal{A} is a subaction of $\text{Init} \wedge \square[\mathcal{N}]_v$.

PROPOSITION 1. *If P is a safety property and L is the conjunction of a countable number of formulas of the form $\text{WF}_w(\mathcal{A})$ and/or $\text{SF}_w(\mathcal{A})$ such that $\mathcal{A} \wedge (w' \neq w)$ is a subaction of P , then (P, L) is machine closed.*

3.4.3 Closure and Hiding. Several of our results have hypotheses of the form $\mathcal{C}(M_1) \wedge \dots \wedge \mathcal{C}(M_n) \Rightarrow \mathcal{C}(M)$. The obvious first step in proving such a formula is to compute the closures $\mathcal{C}(M_1), \dots, \mathcal{C}(M_n)$, and $\mathcal{C}(M)$. We can use Proposition 1 to compute the closure of a formula with no internal variables. When there are internal variables, the following proposition allows us to reduce the proof of $\mathcal{C}(M_1) \wedge \dots \wedge \mathcal{C}(M_n) \Rightarrow \mathcal{C}(M)$ to the proof of a formula in which the closures can be computed with Proposition 1.

PROPOSITION 2. *Let x, x_1, \dots, x_n be tuples of variables such that for each i , no variable in x_i occurs in M or in any M_j with $i \neq j$.*

$$\text{If } \models \bigwedge_{i=1}^n \mathcal{C}(M_i) \Rightarrow \exists x : \mathcal{C}(M), \text{ then } \models \bigwedge_{i=1}^n \mathcal{C}(\exists x_i : M_i) \Rightarrow \mathcal{C}(\exists x : M).$$

Proofs are in the appendix.

Some of our results also have hypotheses of the form $\mathcal{C}(M_1) \wedge \dots \wedge \mathcal{C}(M_n) \Rightarrow E$, where we expect E to be a safety property. If we can verify that E is a safety property, so $\models E = \mathcal{C}(E)$, then we can apply Proposition 2. When E has internal variables, we can often use Proposition 2 of [Abadi and Lamport 1991] to verify that E is a safety property.

3.5 Additional Temporal Operators

We now define some additional temporal operators. Although they can be expressed in terms of the primitive TLA operations $', \square$, and \exists , we define them semantically.

3.5.1 The $+$ Operator. The formula E_{+v} asserts that, if the temporal formula E ever becomes false, then the state function v stops changing. More precisely, a behavior σ satisfies E_{+v} iff either σ satisfies E , or there is some n such that (1) E holds for the first n states of σ and (2) v never changes from the $(n+1)$ st state on. When E is a safety property in canonical form, it is easy to write E_{+v} explicitly:

PROPOSITION 3. *If x is a tuple of variables none of which occurs in v , and s is a variable that does not occur in Init , \mathcal{N} , w , v , or x , and*

$$\begin{aligned}\widehat{\text{Init}} &\triangleq (\text{Init} \wedge (s = 0)) \vee (\neg \text{Init} \wedge (s = 1)) \\ \widehat{\mathcal{N}} &\triangleq \vee(s = 0) \wedge \vee(s' = 0) \wedge (\mathcal{N} \vee (w' = w)) \\ &\quad \vee(s' = 1) \wedge \neg(\mathcal{N} \vee (w' = w)) \\ &\quad \vee(s = 1) \wedge (s' = 1) \wedge (v' = v)\end{aligned}$$

$$\text{then } \models (\exists x : \text{Init} \wedge \square[\mathcal{N}]_w)_{+v} = \exists x, s : \widehat{\text{Init}} \wedge \square[\widehat{\mathcal{N}}]_{(w, v, s)}.$$

We need to reason about $+$ only to verify hypotheses of the form $\models \mathcal{C}(E)_{+v} \wedge \mathcal{C}(M^l) \Rightarrow \mathcal{C}(M)$ in our Decomposition and Composition Theorems. We can verify such a hypothesis by first applying the observation that $\mathcal{C}(E)_{+v}$ equals $\mathcal{C}(E_{+v})$ and using Proposition 3 to calculate E_{+v} . However, this approach is necessary only for noninterleaving specifications. Proposition 4 below provides a way of proving these hypotheses for interleaving specifications without having to calculate E_{+v} .

3.5.2 The \rightarrow Operator. For temporal formulas E and M , the formula $E \rightarrow M$ asserts that M holds at least as long as E does [Abadi and Plotkin 1993]. More precisely $E \rightarrow M$ is true of a behavior σ iff $E \Rightarrow M$ is true of σ and, for every finite prefix ρ of σ , if E is true of ρ then M is true of ρ . It follows from this definition of \rightarrow that $E \rightarrow M$ equals $(\mathcal{C}(E) \rightarrow \mathcal{C}(M)) \wedge (E \Rightarrow M)$. The operator \rightarrow acts much like ordinary implication. In fact, $\models E \rightarrow M$ is equivalent to $\models E \Rightarrow M$. Of course, it is not in general true that $\models (E \rightarrow M) = (E \Rightarrow M)$.

3.5.3 The $\stackrel{+}{\rightarrow}$ Operator. As we observed in the introduction, we interpret the specification that M is guaranteed under assumption E as the formula $E \stackrel{+}{\rightarrow} M$, which means that M holds at least one step longer than E does. More precisely, $E \stackrel{+}{\rightarrow} M$ is true of a behavior σ iff $E \Rightarrow M$ is true of σ and, for every $n \geq 0$, if E holds for the first n states of σ , then M holds for the first $n+1$ states of σ . It follows from this definition of $\stackrel{+}{\rightarrow}$ that $E \stackrel{+}{\rightarrow} M$ equals $(\mathcal{C}(E) \stackrel{+}{\rightarrow} \mathcal{C}(M)) \wedge (E \Rightarrow M)$.

The formula $E \stackrel{+}{\rightarrow} M$ is stronger than $E \rightarrow M$, which asserts that M holds as long as E does. It can be shown that, if E is a safety property, then $E \stackrel{+}{\rightarrow} M$ equals $(M \rightarrow E) \rightarrow M$. We prove in the appendix that, if E and M are both safety properties and v is a tuple of variables containing all free variables of M , then $E \stackrel{+}{\rightarrow} M$ equals $E_{+v} \rightarrow M$.

3.5.4 The \perp Operator. The specification M of a component can be made false only by a step that changes the component's output variables. In an interleaving representation, we do not allow a single step to change output variables of two different components. Hence, if E and M are specifications of separate components, we expect that no step will make both E and M false. More precisely, we expect E and M to be orthogonal (\perp), where $E \perp M$ is true of a behavior σ iff there is no $n \geq 0$ such that E and M are both true for the first n states of σ and both false for the first $n+1$ states of σ . It can be shown that $E \perp M$ equals $\mathcal{C}(E) \perp \mathcal{C}(M)$, and that, if E and M are safety properties, then $E \perp M$ equals $(E \wedge M) \stackrel{+}{\rightarrow} (E \vee M)$.

If no step falsifies both E and M , and M remains true as long as E does, then M must remain true at least one step longer than E does. Hence, $E \perp M$ implies the equivalence of $E \rightarrow M$ and $E \stackrel{+}{\rightarrow} M$. In fact, we prove in the appendix that

$(E \dashv\Rightarrow M) = (E \rightarrow M) \wedge (E \perp M)$ is valid. From this and the relation between $\dashv\Rightarrow$ and $+$, we can derive:

PROPOSITION 4. *If E , M , and R are safety properties, and v is a tuple of variables containing all variables that occur free in M , then $\models E \wedge R \Rightarrow M$ and $\models R \Rightarrow E \perp M$ imply $\models E_{+v} \wedge R \Rightarrow M$.*

This proposition enables us to use orthogonality to remove $+$ from proof obligations. To apply the proposition, we must prove the orthogonality of component specifications. We do this for interleaving specifications with the following result.

PROPOSITION 5. *If*

$$\begin{aligned}\models C(E) &= \text{Init}_E \wedge \square[\mathcal{N}_E]_{(x, e)} \\ \models C(M) &= \text{Init}_M \wedge \square[\mathcal{N}_M]_{(y, m)}\end{aligned}$$

then

$$\models (\exists x : \text{Init}_E \vee \exists y : \text{Init}_M) \wedge \text{Disjoint}(e, m) \Rightarrow C(\exists x : E) \perp C(\exists y : M)$$

4. DECOMPOSING A COMPLETE SPECIFICATION

4.1 Specifying a Component

Let us consider how to write the specification M of one component of a larger system. We assume that the free variables of the specification can be partitioned into tuples m of output variables and e of input variables; the component changes the values of the variables of m only. (A more general situation is discussed below.) The specification of a component has the same form $\exists x : \text{Init} \wedge \square[\mathcal{N}]_v \wedge L$ as that of a complete system. For a component specification:

- v is the tuple $\langle x, m, e \rangle$.
- Init describes the initial values of the component's output variables m and internal variables x .
- \mathcal{N} should allow two kinds of steps—ones that the component performs and ones that its environment performs. Steps performed by the component, which change its output variables m , are described by an action \mathcal{N}_m . In an interleaving representation, the component's inputs and outputs cannot change simultaneously, so \mathcal{N}_m implies $e' = e$. In a noninterleaving representation, \mathcal{N}_m does not constrain the value of e' , so the variables of e do not appear primed in \mathcal{N}_m . In either case, we are specifying the component but not its environment, so we let the environment do anything except change the component's output variables or internal variables. In other words, the environment is allowed to perform any step in which $\langle m, x' \rangle$ equals $\langle m, x \rangle$. (Below, we describe more general specifications in which an environment action can change x .) Therefore, \mathcal{N} should equal $\mathcal{N}_m \vee (\langle m, x' \rangle = \langle m, x \rangle)$.
- L is the conjunction of fairness conditions, each of the form $\text{WF}_{\langle m, x \rangle}(\mathcal{A})$ or $\text{SF}_{\langle m, x \rangle}(\mathcal{A})$. For an interleaving representation, which by definition does not allow steps that change both e and m , the subscripts $\langle m, x \rangle$ and $\langle e, m, x \rangle$ yield equivalent fairness conditions.

This leads us to write M in the form

$$M \triangleq \exists x : \text{Init} \wedge \square[\mathcal{N}_m \vee (\langle m, x \rangle' = \langle m, x \rangle)]_{\langle e, m, x \rangle} \wedge L \quad (3)$$

By simple logic, (3) is equivalent to

$$M \triangleq \exists x : \text{Init} \wedge \square[\mathcal{N}_m]_{\langle m, x \rangle} \wedge L \quad (4)$$

For the specification M_a of process Π_a in the GCD example, x is the empty tuple (there is no internal variable), the input variable e is b , the output variable m is a , and

$$\begin{aligned} \text{Init}_a &\triangleq a = 233344 \\ \mathcal{N}_a &\triangleq (a > b) \wedge (a' = a - b) \wedge (b' = b) \\ M_a &\triangleq \text{Init}_a \wedge \square[\mathcal{N}_a]_a \wedge \text{WF}_a(\mathcal{N}_a) \end{aligned} \quad (5)$$

For the specification M_a^l of the low-level process Π_a^l , the tuple x is $\langle ai, pca \rangle$, where pca is an internal variable that tells whether control is at the beginning of the loop or after the assignment to ai . The specification has the form

$$M_a^l \triangleq \exists ai, pca : \text{Init}_a^l \wedge \square[\mathcal{N}_a^l]_{\langle a, ai, pca \rangle} \wedge \text{WF}_{\langle a, ai, pca \rangle}(\mathcal{N}_a^l) \quad (6)$$

for appropriate initial condition Init_a^l and next-state action \mathcal{N}_a^l . The specifications M_b and M_b^l are similar.

In our queue example, we can write the specifications of both the queue and its environment as separate components in the form (4). For the queue component, the tuple m of output variables is $\langle i.\text{ack}, o.\text{snd} \rangle$, the tuple e of input variables is $\langle i.\text{snd}, o.\text{ack} \rangle$, and the specification is

$$\begin{aligned} IQM &\triangleq \text{Init}_M \wedge \square[\mathcal{Q}_M]_{\langle i.\text{ack}, o.\text{snd}, q \rangle} \wedge ICL \\ QM &\triangleq \exists q : IQM \end{aligned} \quad (7)$$

The specification of the environment as a separate component is

$$QE \triangleq \text{Init}_E \wedge \square[\mathcal{Q}_E]_{\langle i.\text{snd}, o.\text{ack} \rangle} \quad (8)$$

We have provided specifications of the queue and its environment in an interleaving representation. A noninterleaving representation of the queue can be obtained by modifying its specification as follows.

- Change the Enq and Deq actions so they do not constrain the values of $i.\text{snd}'$ or $o.\text{ack}'$.
- Define an action $DeqEng$ that simultaneously enqueues an input value and dequeues an output value, and change the definition of \mathcal{Q}_M to have $DeqEng$ as an additional disjunct.

The resulting specification QM^{ni} is given in Figure 9. It is a noninterleaving specification because it allows a step that changes i and o simultaneously. A noninterleaving representation of the queue's environment can be obtained in a similar fashion.

In describing the component's next-state action \mathcal{N} , we required that an environment action not change the component's internal variables. One can also write a

$$\begin{aligned}
Init_M &\triangleq CInit(o) \wedge (q = \langle \rangle) \\
Enq^{ni} &\triangleq \wedge |q| < N \\
&\quad \wedge Ack(i) \wedge (q' = q \circ \langle i.val \rangle) \\
&\quad \wedge o.snd' = o.snd \\
Deq^{ni} &\triangleq \wedge |q| > 0 \\
&\quad \wedge Send(Head(q), o) \wedge (q' = Tail(q)) \\
&\quad \wedge i.ack' = i.ack \\
DeqEnq^{ni} &\triangleq \wedge (|q| > 0) \wedge Send(Head(q), o) \\
&\quad \wedge Ack(i) \\
&\quad \wedge q' = Tail(q) \circ \langle i.val \rangle \\
Q_M^{ni} &\triangleq Enq^{ni} \vee Deq^{ni} \vee DeqEnq^{ni} \\
IQM^{ni} &\triangleq Init_M \wedge \square [Q_M^{ni}]_{\langle i.ack, o.snd, q \rangle} \wedge WF_{\langle i.ack, o.snd, q \rangle}(Q_M^{ni}) \\
QM^{ni} &\triangleq \exists q : IQM^{ni}
\end{aligned}$$

Fig. 9. A noninterleaving representation of the queue component.

specification in which the component records environment actions by changing its own internal variables. In this case, \mathcal{N} will not equal $\mathcal{N}_m \vee (\langle m, x \rangle' = \langle m, x \rangle)$, but may just imply $(e' = e) \vee (m' = m)$. The resulting formula will not be a pure interleaving specification because environment actions can change the component's variables, but no action can change both the component's and the environment's output variables. We have not explored this style of specification.

We have been assuming that the visible variables of the component's specification can be partitioned into tuples m of output variables and e of input variables. To see how to handle a more general case, let μ_M be the action $m' \neq m$, let v equal $\langle e, m \rangle$, and observe that $[\mathcal{N}_M]_{\langle m, x \rangle}$ equals $[\mathcal{N}_M \vee (\neg \mu_M \wedge (x' = x))]_{\langle v, x \rangle}$. A μ_M step is one that is attributed to the component, since it changes the component's output variables. When the tuple v of variables is not partitioned into input and output variables, we define an action μ_M that specifies what steps are attributed to the component, and we write the component's next-state action in the form $\mathcal{N}_M \vee (\neg \mu_M \wedge (x' = x))$. All our results for separate input and output variables can be generalized by writing the next-state action in this form. However, for simplicity, we consider only the special case.

4.2 Conjoining Components to Form a Complete System

In Section 3.1, we describe how to specify a complete system. In Section 4.1, we describe how to specify an individual component of a system. A complete system is the composition of its components. Composing two systems means constructing a universe in which they are both running. If formulas M_1 and M_2 represent the two systems, then $M_1 \wedge M_2$ represents their composition, since a behavior represents a possible history of a universe containing both systems iff it satisfies both M_1 and M_2 . Thus, in principle, composition is conjunction. We now show that composition is conjunction in practice as well.

For composition to be conjunction, the conjunction of the specifications of all components should be equivalent to the specification of the complete system. For example, the conjunction of the specifications QM of the queue and QE of its environment should be equivalent to the specification CQ of the complete system

shown in Figure 5. Recall that

$$\begin{aligned} QE &= \text{Init}_E \wedge \square[\mathcal{Q}_E]_{\langle i.\text{snd}, o.\text{ack} \rangle} \\ QM &= \exists q : \text{Init}_M \wedge \square[\mathcal{Q}_M]_{\langle i.\text{ack}, o.\text{snd}, q \rangle} \wedge ICL \\ CQ &= \exists q : \wedge \text{Init}_E \wedge \text{Init}_M \\ &\quad \wedge \square \left[\begin{array}{l} \vee \mathcal{Q}_E \wedge (q' = q) \\ \vee \mathcal{Q}_M \end{array} \right]_{\langle i, o, q \rangle} \\ &\quad \wedge ICL \end{aligned}$$

We deduce the equivalence of $QE \wedge QM$ and CQ from the following result, by substituting QE for M_1 and QM for M_2 . (In this case, x_1 is the empty tuple $\langle \rangle$, so \hat{x}_2 equals $\langle \rangle$ and $\hat{x}'_2 = \hat{x}_2$ equals true.)

PROPOSITION 6. *Let $m_1, \dots, m_n, x_1, \dots, x_n$ be tuples of variables, and let*

$$\begin{aligned} m &\triangleq \langle m_1, \dots, m_n \rangle & x &\triangleq \langle x_1, \dots, x_n \rangle \\ \hat{x}_i &\triangleq \langle x_1, \dots, x_{i-1}, x_{i+1}, \dots, x_n \rangle \\ M_i &\triangleq \exists x_i : \text{Init}_i \wedge \square[\mathcal{N}_i]_{\langle m_i, x_i \rangle} \wedge L_i \end{aligned}$$

If, for all $i, j = 1, \dots, n$ with $i \neq j$:

- (1) no variable of x_j occurs free in x_i or M_i ,
- (2) m includes all free variables of M_i , and
- (3) $\models \mathcal{N}_i \Rightarrow (m'_j = m_j)$

then

$$\models \bigwedge_{i=1}^n M_i = \exists x : \bigwedge_{i=1}^n \text{Init}_i \wedge \square[\bigvee_{i=1}^n \mathcal{N}_i \wedge (\hat{x}'_i = \hat{x}_i)]_{\langle m, x \rangle} \wedge \bigwedge_{i=1}^n L_i$$

In this proposition, the third hypothesis asserts that component i leaves the variables of other components unchanged, so M_i is an interleaving representation of component i . Hence, M_i implies $\text{Disjoint}(m_i, m_j)$, for each $j \neq i$, and $\bigwedge_{i=1}^n M_i$ implies $\text{Disjoint}(m_1, \dots, m_n)$, as expected for an interleaving representation of the complete system.

In the GCD example, we apply this proposition to the formula M_a of (5) and the analogous formula M_b . We immediately get that $M_a \wedge M_b$ is equivalent to a formula that is the same as M_{gcd} , defined by (2), except with $\text{WF}_{\langle a, b \rangle}(\mathcal{N}_a) \wedge \text{WF}_{\langle a, b \rangle}(\mathcal{N}_b)$ instead of $\text{WF}_{\langle a, b \rangle}(\mathcal{N})$. It can be shown that these two fairness conditions are equivalent; hence $M_a \wedge M_b$ is equivalent to M_{gcd} .

For another example of decomposition, we consider the system of Figure 7, which consists of two queues in series together with an environment. This system can be decomposed into three components with the following specifications.

$$\begin{aligned} \text{1st queue: } & \exists q_1 : \text{Init}_M^{[1]} \wedge \square[\mathcal{Q}_M^{[1]} \wedge (o' = o)]_{\langle i.\text{ack}, z.\text{snd}, q_1 \rangle} \wedge ICL^{[1]} \\ \text{2nd queue: } & \exists q_2 : \text{Init}_M^{[2]} \wedge \square[\mathcal{Q}_M^{[2]} \wedge (i' = i)]_{\langle z.\text{ack}, o.\text{snd}, q_2 \rangle} \wedge ICL^{[2]} \\ \text{environment: } & \text{Init}_E \wedge \square[\mathcal{Q}_E \wedge (z' = z)]_{\langle i.\text{snd}, o.\text{ack} \rangle} \end{aligned}$$

To obtain an interleaving representation, we have conjoined $o' = o$ to $\mathcal{Q}_M^{[1]}$ in the

ACM Transactions on Programming Languages and Systems, Vol. 17, No. 3, May 1995

first queue's next-state action, because $\mathcal{Q}_M^{[1]}$ does not mention o . Similarly, we have conjoined $i' = i$ to the second queue's next-state action, and $z' = z$ to the environment's. It follows from Proposition 6 that the conjunction of these three specifications equals the specification CDQ of the complete system, defined in Figure 8.

The third hypothesis of Proposition 6 is satisfied only by interleaving representations. For arbitrary representations, a straightforward calculation shows

$$\models \bigwedge_{i=1}^n M_i = \exists x : \wedge \bigwedge_{i=1}^n \text{Init}_i \\ \wedge \square [\bigwedge_{i=1}^n (\mathcal{N}_i \vee \langle m_i, x_i \rangle' = \langle m_i, x_i \rangle)]_{\langle m, x \rangle} \\ \wedge \bigwedge_{i=1}^n L_i \quad (9)$$

assuming only the first hypothesis of the proposition. The right-hand side has the expected form for a noninterleaving specification, since it allows $\mathcal{N}_i \wedge \mathcal{N}_j$ steps for $i \neq j$. Hence, composition is conjunction for noninterleaving representations too.

4.3 The Decomposition Theorem

4.3.1 The Basic Theorem. Consider a complete system decomposed into components Π_i . We would like to prove that this system is implemented by a lower-level one, consisting of components Π_i^l , by proving that each Π_i^l implements Π_i . Let M_i be the specification of Π_i and M_i^l be the specification of Π_i^l . We must prove that $\bigwedge_{i=1}^n M_i^l$ implies $\bigwedge_{i=1}^n M_i$. This implication is trivially true if M_i^l implies M_i , for all i . However, as we saw in the GCD example, M_i^l need not imply M_i .

Even when $M_i^l \Rightarrow M_i$ does not hold, we need not reason about all the lower-level components together. Instead, we prove $E_i \wedge M_i^l \Rightarrow M_i$, where E_i includes just the properties of the other components assumed by component i , and is usually much simpler than $\bigwedge_{k \neq i} M_k^l$. Proving $E_i \wedge M_i^l \Rightarrow M_i$ involves reasoning only about component i , not about the entire lower-level system.

In propositional logic, to deduce that $\bigwedge_{i=1}^n M_i^l$ implies $\bigwedge_{i=1}^n M_i$ from $\bigwedge_{i=1}^n (E_i \wedge M_i^l \Rightarrow M_i)$, we may prove that $\bigwedge_{k=1}^n M_k^l$ implies E_i for each i . However, proving this still requires reasoning about $\bigwedge_{k=1}^n M_k^l$, the specification of the entire lower-level system. The following theorem shows that we need only prove that E_i is implied by $\bigwedge_{k=1}^n M_k$, the specification of the higher-level system—a formula usually much simpler than $\bigwedge_{k=1}^n M_k^l$.

Proving $E_i \wedge M_i^l \Rightarrow M_i$ and $(\bigwedge_{k=1}^n M_k) \Rightarrow E_i$ for each i and deducing $(\bigwedge_{i=1}^n M_i^l) \Rightarrow (\bigwedge_{i=1}^n M_i)$ is circular reasoning, and is not sound in general. Such reasoning would allow us to deduce $(\bigwedge_{i=1}^n M_i^l) \Rightarrow (\bigwedge_{i=1}^n M_i)$ for any M_i^l and M_i —simply let E_i equal M_i . To break the circularity, we need to add some C 's and one hypothesis: if E_i is ever violated then, for at least one additional step, M_i^l implies M_i . This hypothesis is expressed formally as $\models C(E_i)_{+v} \wedge C(M_i^l) \Rightarrow C(M_i)$, for some v ; the hypothesis is weakest when v is taken to be the tuple of all relevant variables. Our proof rule is:

THEOREM 1 (DECOMPOSITION THEOREM). *If, for $i = 1, \dots, n$,*

- (1) $\models \bigwedge_{j=1}^n C(M_j) \Rightarrow E_i$
- (2) (a) $\models C(E_i)_{+v} \wedge C(M_i^l) \Rightarrow C(M_i)$

$$(b) \models E_i \wedge M_i^l \Rightarrow M_i \\ \text{then } \models \bigwedge_{i=1}^n M_i^l \Rightarrow \bigwedge_{i=1}^n M_i.$$

This theorem is a corollary of the Composition Theorem of Section 5.2 below.

In the GCD example, we want to use the theorem to prove $M_a^l \wedge M_b^l \Rightarrow M_a \wedge M_b$. (The component specifications are described in Section 4.1.) The abstract environment specification E_a asserts that b can change only when $a < b$, and that a is not changed by steps that change b . Thus,

$$E_a \triangleq \square[(a < b) \wedge (a' = a)]_b$$

The definition of E_b is analogous. We let v be $\langle a, b \rangle$.

In general, the environment and component specifications can have internal variables. The theorem also allows them to contain fairness conditions. However, the first hypothesis asserts that the E_i are implied by safety properties. In practice, this means that the theorem can be applied only when the E_i are safety properties. The examples of Section 2.2 lead us to expect such a restriction. Moreover, if the E_i have internal variables, we expect them to be simple history-determined variables [Abadi and Lamport 1994], so Proposition 2 of [Abadi and Lamport 1991] can be used to prove that the E_i are safety properties.

4.3.2 Verifying the Hypotheses. We now discuss how one verifies the hypotheses of the Decomposition Theorem, illustrating the method with the GCD example.

To prove the first hypothesis, one first eliminates the closure operators and existential quantifiers by using Propositions 1 and 2 and Proposition 2 of [Abadi and Lamport 1991]. This reduces the hypothesis to a condition of the form

$$\models \bigwedge_{i=1}^n (Init_i \wedge \square[\mathcal{N}_i]_{v_i}) \Rightarrow E_i \quad (10)$$

For interleaving representations, we can then use Proposition 6 to write $\bigwedge_{i=1}^n (Init_i \wedge \square[\mathcal{N}_i]_{v_i})$ in canonical form. For noninterleaving representations, we apply (9). In either case, the proof of (10) is an implementation proof of the kind discussed in Section 3.2.

For the GCD example, the first hypothesis asserts that $C(M_a) \wedge C(M_b)$ implies E_a and E_b . This differs from the third hypothesis of (1) in Section 2.1 because of the C 's. To verify the hypothesis, we can apply Proposition 1 to show that $C(M_a)$ and $C(M_b)$ are obtained by simply deleting the fairness conditions from M_a and M_b . Since \mathcal{N}_b implies $(a < b) \wedge (a' = a)$, it is easy to see that $C(M_b)$ implies E_a . It is equally easy to see that $C(M_a)$ implies E_b . (In more complicated examples, E_i will not follow from $C(M_j)$ for any single j .)

To prove part (a) of the second hypothesis, we first eliminate the $+$. For noninterleaving representations, this must be done with Proposition 3, as described in Section 3.5.1. For interleaving representations, we can apply Propositions 4 and 5, as described in Section 3.5.4. In either case, we can prove the resulting formula by first using Proposition 2 to eliminate quantifiers, using Proposition 1 to compute closures, and then performing a standard implementation proof with a refinement mapping.

Part (b) of the hypothesis also calls for a standard implementation proof, for which we use the same refinement mapping as in the proof of (a). Since E_i implies $\mathcal{C}(E_i)_{+v}$ and M_i^l implies $\mathcal{C}(M_i^l)$, we can infer from part (a) that $E_i \wedge M_i^l$ implies $\mathcal{C}(M_i)$. Thus proving part (b) requires verifying only the liveness part of M_i .

For the GCD example, we verify the two parts of the second hypothesis by proving $\mathcal{C}(E_a)_{+(a,b)} \wedge \mathcal{C}(M_a^l) \Rightarrow \mathcal{C}(M_a)$ and $E_a \wedge M_a^l \Rightarrow M_a$; the proofs of the corresponding conditions for M_b are similar. We first observe that the initial condition of E_a is true, and that, since M_a^l is an interleaving representation, its next-state action \mathcal{N}_a^l implies that no step changes both a and b , so $\mathcal{C}(M_a^l)$ implies $\text{Disjoint}(a, b)$. Hence, applying Propositions 4 and 5, we reduce our task to proving $\mathcal{C}(E_a) \wedge \mathcal{C}(M_a^l) \Rightarrow \mathcal{C}(M_a)$ and $E_a \wedge M_a^l \Rightarrow M_a$. Applying Proposition 2 to remove the quantifier from $\mathcal{C}(M_a^l)$ and Proposition 1 to remove the \mathcal{C} 's, we reduce proving $\mathcal{C}(E_a) \wedge \mathcal{C}(M_a^l) \Rightarrow \mathcal{C}(M_a)$ to proving

$$E_a \wedge \text{Init}_a^l \wedge \square[\mathcal{N}_a^l]_{(a, ai, pca)} \Rightarrow \text{Init}_a \wedge \square[\mathcal{N}_a]_a \quad (11)$$

Using simple logic and (11), we reduce proving $E_a \wedge M_a^l \Rightarrow M_a$ to proving

$$E_a \wedge \text{Init}_a^l \wedge \square[\mathcal{N}_a^l]_{(a, ai, pca)} \wedge \text{WF}_{(a, ai, pca)}(\mathcal{N}_a^l) \Rightarrow \text{WF}_a(\mathcal{N}_a) \quad (12)$$

We can use Proposition 6 to rewrite the left-hand sides of (11) and (12) in canonical form. The resulting conditions are in the usual form for a TLA implementation proof.

In summary, by applying our propositions in a standard sequence, we can use the Decomposition Theorem to reduce decompositional reasoning to ordinary TLA reasoning. This reduction may seem complicated for so trivial an example as the GCD program. However, it will be insignificant compared to the complexity of the complete proof in any realistic example, such as the one by Kurshan and Lamport [1993], discussed below.

4.3.3 The General Theorem. We sometimes need to prove the correctness of systems defined inductively. At induction stage $N+1$, the low- and high-level specifications are defined as the conjunctions of k copies of low- and high-level specifications of stage N , respectively. For example, a 2^{N+1} -bit multiplier is sometimes implemented by combining four 2^N -bit multipliers. We want to prove by induction on N that the stage N low-level specification implements the stage N high-level specification. For such a proof, we need a more general decomposition theorem whose conclusion at stage N can be used in proving the hypotheses at stage $N+1$. The appropriate theorem is:

THEOREM 2 (GENERAL DECOMPOSITION THEOREM). *If, for $i = 1, \dots, n$,*

$$(1) \models \mathcal{C}(E) \wedge \bigwedge_{j=1}^n \mathcal{C}(M_j) \Rightarrow E_i$$

$$(2) (a) \models \mathcal{C}(E_i)_{+v} \wedge \mathcal{C}(M_i^l) \Rightarrow \mathcal{C}(M_i) \\ (b) \models E_i \wedge M_i^l \Rightarrow M_i$$

(3) v is a tuple of variables including all the free variables of M_i

then

$$(a) \models \mathcal{C}(E)_{+v} \wedge \bigwedge_{j=1}^n \mathcal{C}(M_j^l) \Rightarrow \bigwedge_{j=1}^n \mathcal{C}(M_j) \text{ and}$$

$$(b) \models E \wedge \bigwedge_{j=1}^n M_j^l \Rightarrow \bigwedge_{j=1}^n M_j.$$

Conclusion (b) of this theorem has the same form as hypothesis 2(b), with M_i^l and M_i replaced with conjunctions. To make the corresponding hypothesis 2(a) follow from conclusion (a), it suffices to prove $\bigwedge_{j=1}^n \mathcal{C}(M_j) \Rightarrow \mathcal{C}(\bigwedge_{j=1}^n M_j)$, since $\mathcal{C}(\bigwedge_{j=1}^n M_j^l) \Rightarrow \bigwedge_{j=1}^n \mathcal{C}(M_j^l)$ is always true.

The General Decomposition Theorem has been applied to the verification of an inductively defined multiplier circuit [Kurshan and Lamport 1993].

It can be shown that both versions of our decomposition theorem provide complete rules for verifying that one composition implies another. However, this result is of no significance. Decomposition can simplify a proof only if the proof can be decomposed, in the sense that each M_i^l implements the corresponding M_i under a simple environment assumption E_i . Our theorems are designed to handle those proofs that can be decomposed.

5. COMPOSING ASSUMPTION/GUARANTEE SPECIFICATIONS

5.1 The Form of an Assumption/Guarantee Specification

An assumption/guarantee specification asserts that a system guarantees M under the assumption that its environment satisfies E . As we saw in Section 2.2, this specification is expressed by the formula $E \xrightarrow{\perp} M$, which means that, for any n , if the environment satisfies E through “time” n , then the system must satisfy M through “time” $n+1$.

Perhaps the most obvious form for an assumption/guarantee specification is $E \Rightarrow M$. The formula $E \Rightarrow M$ is weaker than $E \xrightarrow{\perp} M$, since it allows behaviors in which M is violated before E . However, an implementation could exploit this extra freedom only by predicting in advance that the environment will violate E . A system does not control its environment, so it cannot predict what the environment will do. The specifications $E \Rightarrow M$ and $E \xrightarrow{\perp} M$ therefore allow the same implementations. We take $E \xrightarrow{\perp} M$ to be the form of assumption/guarantee specifications because this form leads to the simpler rules for composition.

As discussed in Section 2.2, composition works best when environment assumptions are safety properties. It can be shown that $E \xrightarrow{\perp} M$ is equivalent to $\mathcal{C}(E) \xrightarrow{\perp} (\mathcal{C}(M) \wedge (E \Rightarrow M))$, so we can in principle convert any assumption/guarantee specification to one whose assumption is a safety property. (A similar observation appears in our earlier work [Abadi and Lamport 1993, Theorem 1].) However, this equivalence is of intellectual interest only. In practice, we write the environment assumption as a safety property and the system’s fairness guarantee as the conjunction of properties $E_L \Rightarrow \text{WF}_v(\mathcal{A})$ and $E_L \Rightarrow \text{SF}_v(\mathcal{A})$, where E_L is an environment fairness assumption. We can apply Proposition 1 to show that the resulting specification is machine closed because, if (P, L) is machine closed and L implies R , then (P, R) is also machine closed [Abadi and Lamport 1994, Proposition 3].

5.2 The Composition Theorem

Suppose we are given n devices, each with an assumption/guarantee specification $E_j \dashv M_j$. To verify that the composition of these devices implements a higher-level assumption/guarantee specification $E \dashv M$, we must prove $\bigwedge_{j=1}^n (E_j \dashv M_j) \Rightarrow (E \dashv M)$. We use the following theorem:

THEOREM 3 (COMPOSITION THEOREM). *If, for $i = 1, \dots, n$,*

- (1) $\models \mathcal{C}(E) \wedge \bigwedge_{j=1}^n \mathcal{C}(M_j) \Rightarrow E_i$
- (2) (a) $\models \mathcal{C}(E)_{+v} \wedge \bigwedge_{j=1}^n \mathcal{C}(M_j) \Rightarrow \mathcal{C}(M)$
- (b) $\models E \wedge \bigwedge_{j=1}^n M_j \Rightarrow M$

then $\models \bigwedge_{j=1}^n (E_j \dashv M_j) \Rightarrow (E \dashv M)$.

This theorem also allows us to prove conditional implementation results of the form $G \wedge \bigwedge_{j=1}^n (E_j \dashv M_j) \Rightarrow (E \dashv M)$; we just let M_1 equal G and E_1 equal true, since true $\dashv G$ equals G . For interleaving specifications, we can in general prove only conditional implementation, where G includes disjointness conditions asserting that the outputs of different components do not change simultaneously.

The hypotheses of the Composition Theorem are similar to those of the Decomposition Theorem, and they are proved in much the same way. The major difference is that, for interleaving specifications, the orthogonality condition $\mathcal{C}(E) \perp \mathcal{C}(M)$ does not follow from the form of the component specifications, but requires explicit disjointness assumptions.

Observe that the hypotheses have the form $\models P \wedge \bigwedge_{j=1}^n Q_j \Rightarrow R$. Each formula $P \wedge \bigwedge_{j=1}^n Q_j$ has the form of the specification of a complete system, with component specifications P, Q_1, \dots, Q_n . Thus, each hypothesis asserts that a complete system satisfies a property R . In other words, the theorem reduces reasoning about assumption/guarantee specifications to the kind of reasoning used for complete-system specifications.

Among the corollaries of the Composition Theorem are ones that allow us to prove that a lower-level specification implies a higher-level one. The simplest such result has, as its conclusion, $\models (E \dashv M^l) \Rightarrow (E \dashv M)$. This condition expresses the correctness of the refinement of a component with a fixed environment assumption.

COROLLARY 1. *If E is a safety property and*

- (a) $\models E_{+v} \wedge \mathcal{C}(M^l) \Rightarrow \mathcal{C}(M)$
 - (b) $\models E \wedge M^l \Rightarrow M$
- then* $\models (E \dashv M^l) \Rightarrow (E \dashv M)$.

5.3 The Queue Example

The assumption/guarantee specification of the queue of Figure 2 is $QE \xrightarrow{\pm} QM$, where QM and QE are defined in (7) and (8) of Section 4.1. We now compose two queues, as shown in Figure 7. The specifications of these queues are obtained from $QE \xrightarrow{\pm} QM$ by substitution; they are $QE^{[1]} \xrightarrow{\pm} QM^{[1]}$ and $QE^{[2]} \xrightarrow{\pm} QM^{[2]}$. We want to show that their composition implements the $(2N+1)$ -element queue specified by $QE^{[\text{dbl}]} \xrightarrow{\pm} QM^{[\text{dbl}]}$. The obvious thing to try to prove is

$$(QE^{[1]} \xrightarrow{\pm} QM^{[1]}) \wedge (QE^{[2]} \xrightarrow{\pm} QM^{[2]}) \Rightarrow (QE^{[\text{dbl}]} \xrightarrow{\pm} QM^{[\text{dbl}]}) \quad (13)$$

We could prove this had we used a noninterleaving representation of the queue. However, (13) is not valid for an interleaving representation, for the following reason. The specification of the first queue does not mention o , and that of the second queue does not mention i . The conjunction of the two specifications allows an enqueue action of the first queue and a dequeue action of the second queue to happen simultaneously, a step that changes $i.\text{ack}$ and $o.\text{snd}$ simultaneously. But, in an interleaving representation, the $(2N+1)$ -element queue's guarantee does not allow such a step, so (13) must be invalid. Another problem with (13) is that the conjunction of the component queues' specifications allows a step that changes $z.\text{snd}$ and $o.\text{ack}$ simultaneously. Such a step satisfies the $(2N+1)$ -element queue's environment assumption $QE^{[\text{dbl}]}$, which does not mention z , so (13) asserts that the next step must satisfy its guarantee $QM^{[\text{dbl}]}$. However, a step that changes both $z.\text{snd}$ and $o.\text{ack}$ violates the second component queue's environment assumption $QE^{[2]}$, permitting the component queue to make arbitrary changes to $o.\text{snd}$ in the next step. A similar problem is caused by simultaneous changes to $i.\text{snd}$ and $z.\text{ack}$.

We already faced the problem of disallowing simultaneous changes to different components' outputs in Section 4.2, where we decomposed an interleaving specification of a $(2N+1)$ -element queue. There, the solution was to strengthen the next-state actions of the component queues and of the environment. This solution cannot be used if we want to compose preexisting specifications without modifying them. In this case, we prove that the composition implements the larger queue under the assumption that the outputs of two different components do not change simultaneously. Thus, we prove

$$G \wedge (QE^{[1]} \xrightarrow{\pm} QM^{[1]}) \wedge (QE^{[2]} \xrightarrow{\pm} QM^{[2]}) \Rightarrow (QE^{[\text{dbl}]} \xrightarrow{\pm} QM^{[\text{dbl}]}) \quad (14)$$

where G is the formula

$$G \triangleq \text{Disjoint}(\langle i.\text{snd}, o.\text{ack} \rangle, \langle z.\text{snd}, i.\text{ack} \rangle, \langle o.\text{snd}, z.\text{ack} \rangle)$$

The proof is outlined in Figure 10.

6. CONCLUSION

We have developed a method for describing components of concurrent systems as TLA formulas. We have shown how to describe a complete system as the conjunction of component specifications and how to describe an open system as a formula $E \xrightarrow{\pm} M$, where E and M are specifications of an environment component and a system component, respectively. Although the idea of reducing programming concepts to logic is old, our approach is new. Our style of writing specifications is direct and, we believe, practical.

$$1. \mathcal{C}(QE^{[\text{dbl}]}) \wedge \mathcal{C}(G) \wedge \mathcal{C}(QM^{[1]}) \wedge \mathcal{C}(QM^{[2]}) \Rightarrow QE^{[1]} \wedge QE^{[2]}$$

PROOF: We use Propositions 2 and 1 to remove the quantifiers and closure operators from the left-hand side of the implication. The resulting formula then asserts that a complete system, consisting of the safety parts of the two queues (with their internal state visible) together with the environment, implements $QE^{[1]}$ and $QE^{[2]}$. The proof of this formula is straightforward.

$$2. \mathcal{C}(QE^{[\text{dbl}]})_{+(i,o,z)} \wedge \mathcal{C}(QM^{[1]}) \wedge \mathcal{C}(G) \wedge \mathcal{C}(QM^{[2]}) \Rightarrow \mathcal{C}(QM^{[\text{dbl}]})$$

$$2.1. \mathcal{C}(G) \wedge \mathcal{C}(QM^{[1]}) \wedge \mathcal{C}(QM^{[2]}) \Rightarrow \mathcal{C}(QE^{[\text{dbl}]}) \perp \mathcal{C}(QM^{[\text{dbl}]})$$

$$2.1.1. \mathcal{C}(IQM^{[1]}) \wedge \mathcal{C}(IQM^{[2]}) \Rightarrow \exists q_1, q_2 : \text{Init}_M^{[1]} \wedge \text{Init}_M^{[2]}$$

PROOF: Follows easily from Proposition 1 and the definitions.

$$2.1.2. \mathcal{C}(QM^{[1]}) \wedge \mathcal{C}(QM^{[2]}) \Rightarrow \exists q_1, q_2 : \text{Init}_M^{[1]} \wedge \text{Init}_M^{[2]}$$

PROOF: 2.1.1 and Proposition 2 (since any predicate is a safety property).

2.1.3. Q.E.D.

PROOF: 2.1.2, the definition of G , and Proposition 5 (since disjointness is a safety property).

$$2.2. \mathcal{C}(QE^{[\text{dbl}]}) \wedge \mathcal{C}(G) \wedge \mathcal{C}(QM^{[1]}) \wedge \mathcal{C}(QM^{[2]}) \Rightarrow \mathcal{C}(QM^{[\text{dbl}]})$$

PROOF: We use Propositions 2 and 1 to remove the quantifiers and closures from the formula. The resulting formula is proved when proving the safety part of step 3.

2.3. Q.E.D.

PROOF: 2.1, 2.2, and Proposition 4.

$$3. QE^{[\text{dbl}]} \wedge G \wedge QM^{[1]} \wedge QM^{[2]} \Rightarrow QM^{[\text{dbl}]}$$

PROOF: A direct calculation shows that the left-hand side of the implication implies CDQ , the complete-system specification of the double queue. We already observed in Section 3.2 that CDQ implements $CQ^{[\text{dbl}]}$, which equals $QE^{[\text{dbl}]} \wedge QM^{[\text{dbl}]}$.

4. Q.E.D.

PROOF: 1–3 and the Composition Theorem, substituting

$$\begin{array}{llll} M_1 \leftarrow G & M_2 \leftarrow QM^{[1]} & M_3 \leftarrow QM^{[2]} & M \leftarrow QM^{[\text{dbl}]} \\ E_1 \leftarrow \text{true} & E_2 \leftarrow QE^{[1]} & E_3 \leftarrow QE^{[2]} & E \leftarrow QE^{[\text{dbl}]}\end{array}$$

Fig. 10. Proof sketch of (14).

We have also provided rules for proving properties of large systems by reasoning about their components. The Composition and Decomposition Theorems are rather simple, yet they allow fairness properties and hiding. They were preceded by results in a long list of publications, described next.

Like ours, most previous composition theorems were strong, in the sense that they could handle circularities for safety properties. Our approach differs from earlier ones in its general treatment of fairness and hiding. The first strong composition theorem we know is that of Misra and Chandy [1981], who considered safety properties of processes communicating by means of CSP primitives. They wrote assumption/guarantee specifications as Hoare triples containing assertions about history variables. Pandya and Joseph [1991] extended this approach to handle some liveness properties. Pnueli [1984] was the first to use temporal logic to write assumption/guarantee specifications. He had a strong composition theorem for safety properties with no hiding. To handle liveness, he wrote assumption/guarantee specifications with implication instead of \rightarrow , so he did not obtain a strong composition theorem. Stark [1985] also wrote assumption/guarantee specifications as implications of temporal formulas and required that circularity be avoided. Our earlier work [Abadi and Lamport 1993] was semantic, in a more complicated model with agents. It lacked practical proof rules for handling fairness and hiding. Collette [1993] adapted this work to Unity. Abadi and Plotkin [1993] used a propositional logic with agents, and considered only safety properties.

Most previous papers were concerned only with composition of assumption/guarantee specifications, and lacked an analog of our Decomposition Theorem. An exception is the work of Berthet and Cerny [1988], who used decomposition in proving safety properties for finite-state automata.

So far, we have applied our Composition Theorem only to toy examples. Formal reasoning about systems is still rare, and it generally occurs on a case-by-case basis. When the specification of a component is used only to verify a specific system, there is no need for a general assumption/guarantee specification. For most practical applications, decomposition suffices. When decomposition does not suffice, the Composition Theorem makes reasoning about open systems almost as easy as reasoning about complete ones.

We have used our Decomposition Theorem with no difficulty on a few toy examples. However, we believe that its biggest payoff will be for systems that are too complex to verify easily by hand. The theorem makes it possible for decision procedures to do most of the work in verifying a system, even when these procedures cannot be applied to the whole system because its state space is very large or unbounded. This approach is currently being pursued in one substantial example: the mechanical verification of a multiplier circuit using a combination of TLA reasoning and mechanical verification with COSPAN [Kurshan and Lamport 1993]. Because it eliminates reasoning about the complete low-level system, the Decomposition Theorem is the key to this division of labor.

Acknowledgments

Karlis Cerans, Stephan Merz, Yuan Yu, and anonymous referees provided helpful comments on an earlier version.

ACM Transactions on Programming Languages and Systems, Vol. 17, No. 3, May 1995.

REFERENCES

- ABADI, M. AND LAMPORT, L. 1994. An old-fashioned recipe for real time. *ACM Trans. Program. Lang. Syst.* 16, 5 (Sept.), 1543–1571.
- ABADI, M. AND LAMPORT, L. 1993. Composing specifications. *ACM Trans. Program. Lang. Syst.* 15, 1 (Jan.), 73–132.
- ABADI, M. AND LAMPORT, L. 1991. The existence of refinement mappings. *Theor. Comput. Sci.* 82, 2 (May), 253–284.
- ABADI, M. AND PLOTKIN, G. 1993. A logical view of composition. *Theor. Comput. Sci.* 114, 1 (June), 3–30.
- ABRAMSKY, S. AND JAGADEESAN, R. 1994. Games and full completeness for multiplicative linear logic. *J. Symb. Logic* 59, 2 (June), 543–574.
- ALPERN, B. AND SCHNEIDER, F. B. 1985. Defining liveness. *Inf. Process. Lett.* 21, 4 (Oct.), 181–185.
- BERTHET, C. AND CERNY, E. 1988. An algebraic model for asynchronous circuits verification. *IEEE Trans. Comput.* 37, 7 (July), 835–847.
- COLLETTE, P. 1993. Application of the composition principle to Unity-like specifications. In *TAPSOFT'93: Theory and Practice of Software Development*, M.-C. Gaudel and J.-P. Jouannaud, Eds. Lecture Notes in Computer Science, vol. 668. Springer-Verlag, Berlin, 230–242.
- JONES, C. B. 1983. Specification and design of (parallel) programs. In *Information Processing 83: Proceedings of the IFIP 9th World Congress*, R. E. A. Mason, Ed. North-Holland, Amsterdam, 321–332.
- KURSHAN, R. P. AND LAMPORT, L. 1993. Verification of a multiplier: 64 bits and beyond. In *Computer-Aided Verification*, C. Courcoubetis, Ed. Lecture Notes in Computer Science, vol. 697. Springer-Verlag, Berlin, 166–179. Proceedings of the 5th International Conference, CAV'93.
- LAMPORT, L. 1994. The temporal logic of actions. *ACM Trans. Program. Lang. Syst.* 16, 3 (May), 872–923.
- LAMPORT, L. 1989. A simple approach to specifying concurrent systems. *Commun. ACM* 32, 1 (Jan.), 32–45.
- LAMPORT, L. 1983. What good is temporal logic? In *Information Processing 83: Proceedings of the IFIP 9th World Congress*, R. E. A. Mason, Ed. North-Holland, Amsterdam, 657–668.
- MEAD, C. AND CONWAY, L. 1980. *Introduction to VLSI Systems*. Addison-Wesley, Reading, Mass., chap. 7.
- MISRA, J. AND CHANDY, K. M. 1981. Proofs of networks of processes. *IEEE Trans. Softw. Eng.* SE-7, 4 (July), 417–426.
- PANDYA, P. K. AND JOSEPH, M. 1991. P-A logic—a compositional proof system for distributed programs. *Distrib. Comput.* 5, 1, 37–54.
- PNUELI, A. 1984. In transition from global to modular temporal reasoning about programs. In *Logics and Models of Concurrent Systems*, K. R. Apt, Ed. NATO ASI Series, Springer-Verlag, Berlin, 123–144.
- PNUELI, A. 1981. The temporal semantics of concurrent programs. *Theor. Comput. Sci.* 13, 45–60.
- STARK, E. W. 1985. A proof technique for rely/guarantee properties. In *Foundations of Software Technology and Theoretical Computer Science*, S. N. Maheshwari, Ed. Lecture Notes in Computer Science, vol. 206. Springer-Verlag, Berlin, 369–391.
- ZAVE, P. AND JACKSON, M. 1993. Conjunction as composition. *ACM Trans. Softw. Eng. Method.* 2, 4 (Oct.), 379–411.

Received December 1993; revised July 1994 and January 1995; accepted January 1995

ACM Transactions on Programming Languages and Systems, Vol. 17, No. 3, May 1995.

